

Unsupervised Learning II

Dr. Alex Williams

October 21, 2020



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE

COSC 425: Introduction to Machine Learning
Fall 2020 (CRN: 44874)

Today's Agenda



We will address:

1. Dimensionality Reduction with PCA

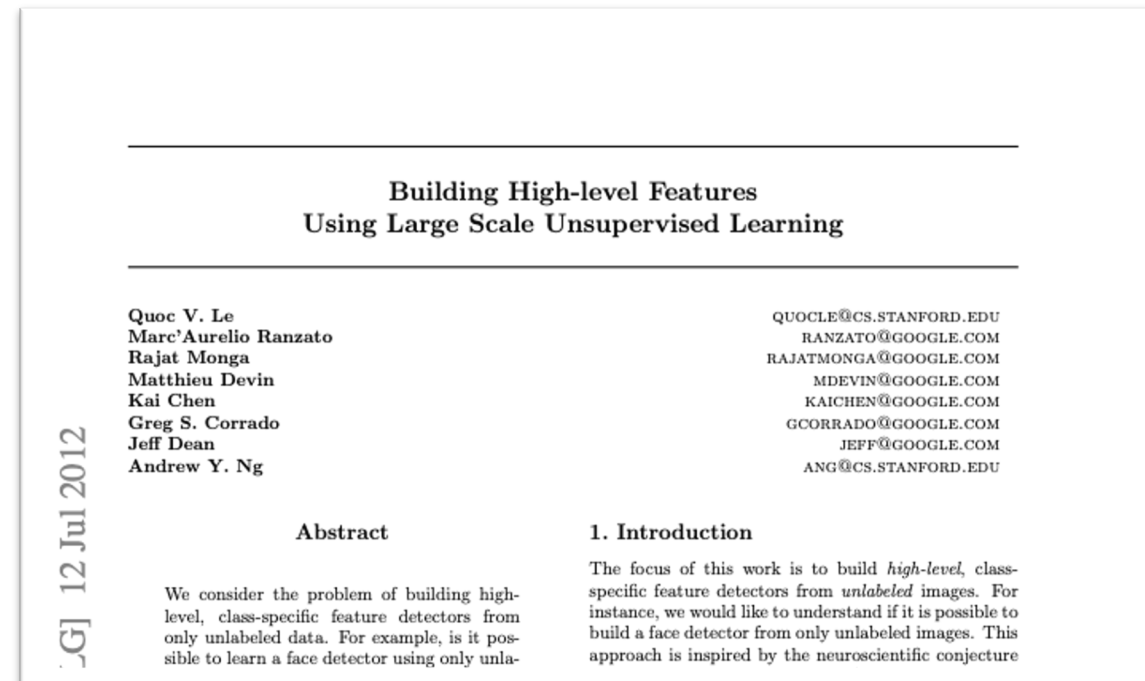
Unsupervised Learning: Feature Selection

Long-Term Goal.

- Figure out which inputs matter.

Feasible, but Challenging.

- Data, data, and more data.



+2000 Citations!

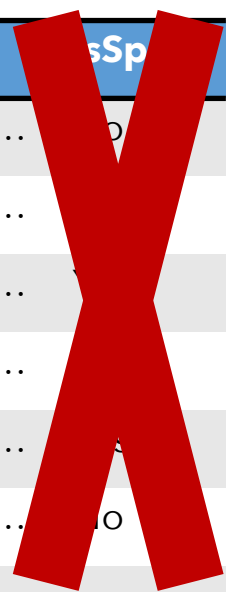
<https://arxiv.org/pdf/1112.6209.pdf>

Supervised Learning: Deterrents

Challenges in Practice

- Finding a "labeler" may be difficult, expensive or impossible.
- Unsupervised learning is concerned with learning without a teacher.

	isUTKEmail	HeaderKeyword	Word 1	Word 2	...	isSpam
x1	Yes	CS425	Hi	Prof	...	No
x2	Yes	Orientation	Alex	You	...	No
x2	No	urgent	Dear	Sir	...	No
x4	No	cash	hello	I	...	No
x5	No	help	are	you	...	No
x6	Yes	Survey	Faculty	this	...	No
		...				



Unsupervised Learning

In unsupervised learning, data consists only of examples and not the corresponding labels.

→ Our job is to make sense of or find some pattern of regularity in the data even though no one has provided correct labels.

For example, we might want to do:



- **Clustering:** Automatically partition the data into groups.
- **Dimensionality Reduction:** Project high dimension data into lower dimension space, so it can be more easily visualized.

Overview

1. Dimensionality Reduction with PCA

Dimensionality Reduction

- **mpg**: continuous
- **cylinders**: multi-valued discrete
- **displacement**: continuous
- **horsepower**: continuous
- **weight**: continuous
- **acceleration**: continuous
- **model year**: multi-valued discrete
- **origin**: multi-valued discrete
- **car name**: string (unique for each instance)



How do we automatically detect and remove redundant dimensions?

We could've done this naively by looking at R-squared metrics in Project 2.

Principal Component Analysis

Dimensionality reduction is the task of taking a dataset in high dimensions (say 10,000) and reducing it to low dimensions (say 2) while retaining the “important” characteristics of the data.

→ Unsupervised learning context = important is difficult to define.

Principal Component Analysis (PCA)

→ Extracts “important” information from a multivariate dataset and expresses the information as a set of new variables

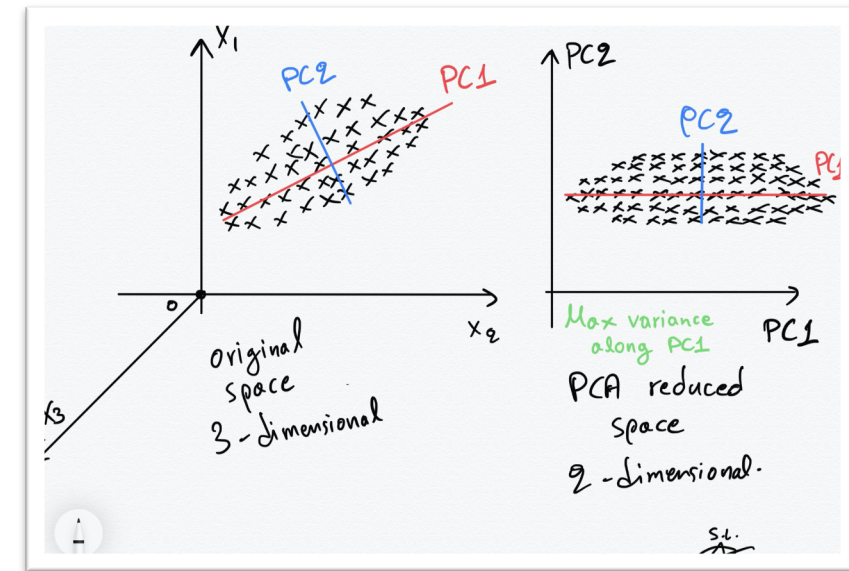
→ **“Principal Components”**

Number is less than or equal to # of features.

→ Allows us to summarize and to visualize information in a dataset of high dimensions.

→ 2D to 1D

→ 3D to 2D



Credit: [Towards Data Science](#)

PCA: Solution + Algorithm

PCA operates by looking for a vector \mathbf{u} in your feature space with the maximal variance.

Procedure

1. Normalize Features (This is a pre-processing step!)
→ Ensure each feature has a zero mean and is scales accordingly.
2. Compute Covariance Matrix
→ Quantifies variance between features.
3. Compute Eigenvectors
→ Vectors that allow us to transform our dimensionality.
4. Keep the first k Eigenvectors and project to get new features \mathbf{z} .
→ Translated into principal components.

Normalization and Scaling

Normalization is PCA's pre-processing step that involves computing the mean of all features to orient around zero.

→ Also known as "Mean Normalization".

→ Each x_j^i is replaced with $x_j - \mu_j$


$$\mu_j = \frac{1}{m} \sum_{i=1}^m x_j^{(i)}$$

Scaling

You should make your features are on the same scale

→ X_1 = Size of House

→ X_2 = Number of Bedrooms



We need to scale features to have a comparable range of values.

→ This is precisely what we did in our Linear Regression assignment.

PCA: Example

Step 1: Apply mean normalization.

		x	y		
Data =		2.5	2.4	x	y
		0.5	0.7	.69	.49
		2.2	2.9	-1.31	-1.21
		1.9	2.2	.39	.99
		3.1	3.0	.09	.29
		2.3	2.7	1.29	1.09
		2	1.6	.49	.79
		1	1.1	.19	-.31
		1.5	1.6	-.81	-.81
		1.1	0.9	-.31	-.31
				-.71	-1.01

$x_i - \bar{x}$
 $y_i - \bar{y}$
 $\forall i = 1, \dots, 10$

→ DataAdjust =

[Clark, 2002. PCA.](#)

http://www.cs.otago.ac.nz/cosc453/student_tutorials/principal_components.pdf

PCA: Example

Step 2: Calculate Covariance Matrix

Covariance measures how two variables vary from the mean with respect to one another.

$$\text{var}(X) = \frac{\sum_{i=1}^n (X_i - \bar{X})(X_i - \bar{X})}{(n - 1)}$$

$$\text{cov}(X) = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{(n - 1)}$$

Covariance matrix captures covariance values between all dimensions.

$$C = \begin{pmatrix} \text{cov}(x^1, x) & \text{cov}(x^2, y) \\ \text{cov}(y^3, x) & \text{cov}(y^4, y) \end{pmatrix}$$

$$C = \begin{pmatrix} 0.6166 & 0.6154 \\ 0.6154 & 0.7166 \end{pmatrix}$$

non-diagonal values are positive, indicating that x increases as y increases.

PCA: Example

Step 3: Calculate the Eigenvectors and Eigenvalues of the Covariance matrix.


Eigenvector v of a linear transformation is a vector that, upon transformation, does not change direction.

$$Av = \lambda v$$

$$\begin{pmatrix} 2 & 3 \\ 2 & 1 \end{pmatrix} \times \begin{pmatrix} 3 \\ 2 \end{pmatrix} = \begin{pmatrix} 12 \\ 8 \end{pmatrix} = 4 \times \begin{pmatrix} 3 \\ 2 \end{pmatrix}$$

Here, the associated eigenvalue is 4.

→ Not that all eigenvectors of a matrix are orthogonal (i.e. perpendicular) to each other. We can re-express the data using eigenvectors as new axes.



Blue Arrow = Eigenvector

- Doesn't change direction
- Eigenvector is 1 because length doesn't change.

PCA: Example

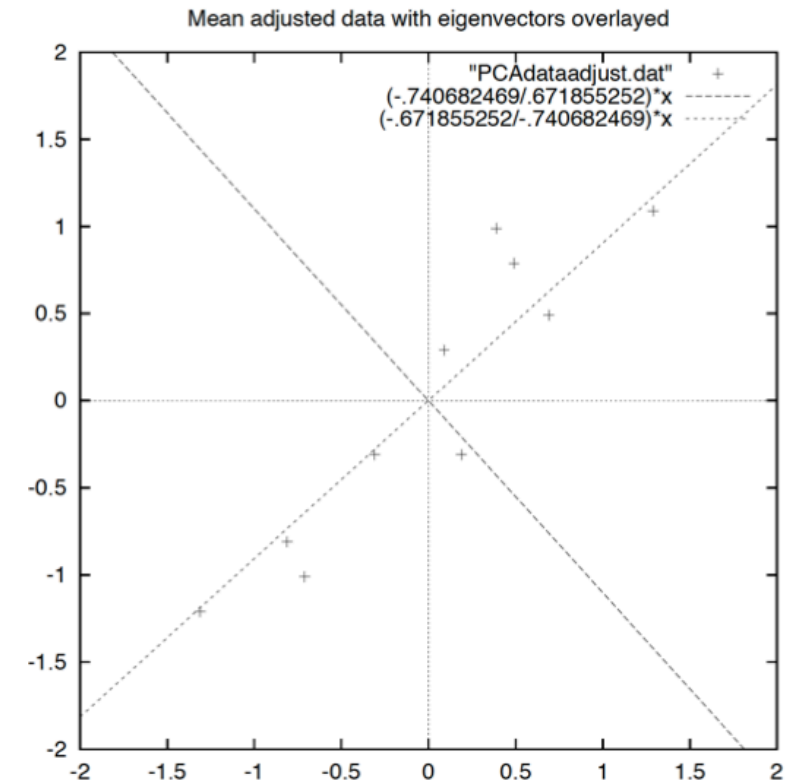
Step 3: Calculate the Eigenvectors and Eigenvalues of the Covariance matrix.

The Unit Eigenvectors

$$\text{eigenvectors} = \begin{pmatrix} -0.7352 & -0.6779 \\ 0.6779 & -0.7352 \end{pmatrix}$$

The Corresponding Eigenvalues

$$\text{eigenvalues} = \begin{pmatrix} 0.049 \\ 1.284 \end{pmatrix}$$



PCA: Example

Step 4: Choose Principal Components

The Eigenvector with the highest eigenvalue is the principal component of the dataset.

$$\text{eigenvectors} = \begin{pmatrix} -0.7352 & -0.6779 \\ 0.6779 & -0.7352 \end{pmatrix} \quad \text{eigenvalues} = \begin{pmatrix} 0.049 \\ 1.284 \end{pmatrix}$$

Form a matrix of k eigenvectors, ordered by eigenvalues from largest to smallest.

$$W = [eig_1, eig_2, \dots, eig_k]$$

$$W = \begin{pmatrix} -0.6779 & -0.7352 \\ -0.7352 & 0.6779 \end{pmatrix}$$

If $k < m$, then you are essentially discarding some dimensions. e.g.,

$$W = \begin{pmatrix} -0.6779 \\ -0.7352 \end{pmatrix}$$

PCA: Example

Step 5: Derive the new Dataset

Multiply the transposition of W on the left of the mean-adjusted dataset, transposed.

Eigenvectors (Step 4)

Mean Normalized Data (Step 1)

$$\begin{pmatrix} -0.6779 & -0.7352 \\ -0.7352 & 0.6779 \end{pmatrix} \begin{pmatrix} 0.69 & -1.31 & 0.39 & 0.09 & 1.29 & 0.49 & 0.19 & -0.81 & -0.31 & -0.71 \\ 0.49 & -1.21 & 0.99 & 0.29 & 1.09 & 0.79 & -0.31 & -0.81 & -0.31 & -1.01 \end{pmatrix}$$


=

x	y
-0.827970186	-0.175115307
1.77758033	.142857227
-0.992197494	.384374989
-0.274210416	.130417207
-1.67580142	-.209498461
-.912949103	.175282444
.0991094375	-.349824698
1.14457216	.0464172582
.438046137	.0177646297
1.22382056	-.162675287

PCA: Example

Step 5: Derive the new Dataset

x	y
-.827970186	-.175115307
1.77758033	.142857227
-.992197494	.384374989
-.274210416	.130417207
-1.67580142	-.209498461
-.912949103	.175282444
.0991094375	-.349824698
1.14457216	.0464172582
.438046137	.0177646297
1.22382056	-.162675287



Transformed Data (Single eigenvector)

x
-.827970186
1.77758033
-.992197494
-.274210416
-1.67580142
-.912949103
.0991094375
1.14457216
.438046137
1.22382056

PCA with Sklearn

```
PCA-Example > 🐞 pca.py > ...  
1  import numpy as np  
2  import pandas as pd  
3  import matplotlib.pyplot as plt  
4  from mpl_toolkits.mplot3d import Axes3D  
5  from sklearn.decomposition import PCA  
6  from sklearn.preprocessing import StandardScaler
```

Versions

- Python 3.7
- Matplotlib=3.1.2
- Numpy =1.17.2
- Pandas=0.25.1
- Scikit-learn=0.23.2

pca.py on Canvas

PCA with Sklearn

```
7  
8 url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"  
9  
10 # load dataset into Pandas DataFrame  
11 df = pd.read_csv(url, names=['sepal length', 'sepal width', 'petal length', 'petal width', 'target'])  
12
```

	sepal length	sepal width	petal length	petal width	target
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

The Iris Dataset

PCA with Sklearn

	sepal length	sepal width	petal length	petal width
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2

Standardization

	sepal length	sepal width	petal length	petal width
0	-0.900681	1.032057	-1.341272	-1.312977
1	-1.143017	-0.124958	-1.341272	-1.312977
2	-1.385353	0.337848	-1.398138	-1.312977
3	-1.506521	0.106445	-1.284407	-1.312977
4	-1.021849	1.263460	-1.341272	-1.312977

```
14 features = ['sepal length', 'sepal width', 'petal length', 'petal width']
15 # Separating out the features
16 x = df.loc[:, features].values
17 # Separating out the target
18 y = df.loc[:, ['target']].values
19 # Standardizing the features
20 x = StandardScaler().fit_transform(x)
```

Standardization

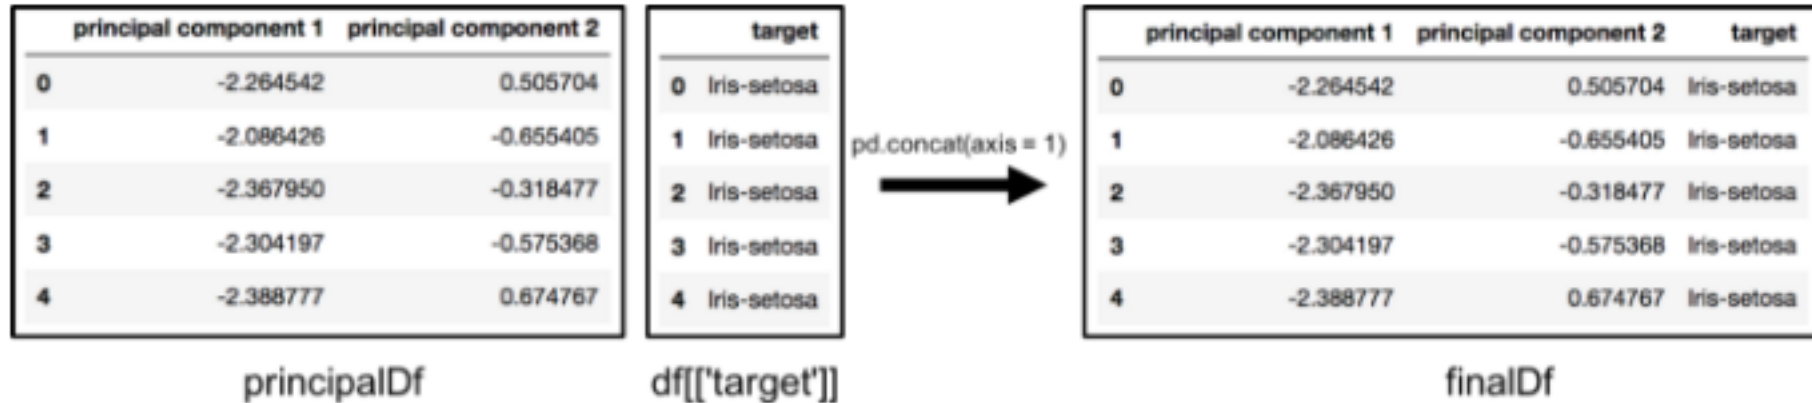
PCA with Sklearn

	sepal length	sepal width	petal length	petal width		principal component 1	principal component 2
0	-0.900681	1.032057	-1.341272	-1.312977	PCA (2 components) →	-2.264542	0.505704
1	-1.143017	-0.124958	-1.341272	-1.312977		-2.086426	-0.655405
2	-1.385353	0.337848	-1.398138	-1.312977		-2.367950	-0.318477
3	-1.506521	0.106445	-1.284407	-1.312977		-2.304197	-0.575368
4	-1.021849	1.263460	-1.341272	-1.312977		-2.388777	0.674767

```
23  pca = PCA(n_components=2)
24  principalComponents = pca.fit_transform(x)
25  principalDf = pd.DataFrame(data = principalComponents
26  |   |   |   | , columns = ['principal component 1', 'principal component 2'])
27
```

Apply PCA: 4D to 2D Space


PCA with Sklearn



```
28 finalDf = pd.concat([principalDf, df[['target']], axis = 1)
```

Concatenate the Target Label into the PC dataframe

SVMs with Scikit-learn




Visualizing Boundaries

Source code included in the svm.py file on Canvas.

PCA with Sklearn

```
30 fig = plt.figure(figsize = (8,8))
31 ax = fig.add_subplot(1,1,1)
32 ax.set_xlabel('Principal Component 1', fontsize = 15)
33 ax.set_ylabel('Principal Component 2', fontsize = 15)
34 ax.set_title('2 component PCA', fontsize = 20)
35 targets = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']
36 colors = ['r', 'g', 'b']
37 for target, color in zip(targets, colors):
38     indicesToKeep = finalDf['target'] == target
39     ax.scatter(finalDf.loc[indicesToKeep, 'principal component 1']
40             , finalDf.loc[indicesToKeep, 'principal component 2']
41             , c = color
42             , s = 50)
43 ax.legend(targets)
44 ax.grid()
45 plt.show()
```

Plot PCs in 2D Space



PCA with Sklearn

```
47 print("How much of our variance is explained?")
48 print(pca.explained_variance_ratio_)
49 print()
50 print()
51
52 print("Which features matter most?")
53 print(abs(pca.components_))
```

```
[(base) alex@MacBook-Pro:~/Desktop/COSC425/PCA-Example$ python pca.py
How much of our variance is explained?
[0.72770452 0.23030523]
Which features matter most?
[[0.52237162 0.26335492 0.58125401 0.56561105]
 [0.37231836 0.92555649 0.02109478 0.06541577]]
```

Explained Variance

- 72.7% for PC1
- 23.0% for PC2

Combined: 95.7%

- PCA loses only 4.3% of the "information" in our data.

Feature Importance

- 1st row = PC1 → Features 1, 3, and 4 matter equally.
- 2nd row = PC2 → Features 2 matters *a lot*!

OK. Now, what?

- Say Feature 3 had low importance in both PC1 and PC2.
 - Remove it from the feature set entirely.

Next Time



We will address:

1. Neural Nets